

Introduction to R

Rutgers


Jason Bryer, Ph.D.
February 14, 2016

Agenda

- Overview
- Loading Data
- The R Language
- Statistics
- Data Visualization
- Document Preparation with R Markdown
- Advanced Statistics and Predictive Modeling
- Interactive Documents with Shiny

Installation

- Instructions for R are here: <https://github.com/jbryer/introR/blob/master/Installation/install.md>
 - R: <https://cran.r-project.org/>
 - RStudio: <https://www.rstudio.com/products/rstudio/download/>
- Download slides here: <https://github.com/jbryer/talks>



What is R?

"R is a language and environment for statistical computing and graphics. It is a GNU project which is similar to the S language and environment which was developed at Bell Laboratories (formerly AT&T, now Lucent Technologies) by John Chambers and colleagues..."

"R provides a wide variety of statistical (linear and non linear modeling, classical statistical tests, time-series analysis, classification, clustering, ...) and graphical techniques, and is highly extensible. The S language is often the vehicle of choice for research in statistical methodology, and R provides an Open Source route to participation in that activity."

(R-project.org)

Overview

Pros

- FREE! *R* is available as Free Software under the terms of the Free Software Foundation's GNU General Public License in source code form. It compiles and runs on a wide variety of UNIX platforms and similar systems (including FreeBSD and Linux), Windows and MacOS.
- Available for multiple platforms (i.e. Windows, Mac, Linux).
- Easily extensible with (currently) over 2,000 packages listed on CRAN.
- Scriptable.
- Publication grade graphics.
- Multiple ways of doing the same thing.
- Quickly becoming the *de facto* standard among statisticians.


7/123

Cons

- Has a steeper learning curve.
- Multiple ways of doing the same thing.
- Can have difficulty with very large datasets.

8/123

The Popularity of R



Firth, D (2011). R and citations. Weblog entry at URL <https://statgeek.wordpress.com/2011/06/25/r-and-citations/>.

See also: Muenchen, R.A. (2017). The Popularity of Data Analysis Software. Weblog entry at URL <http://r4stats.com/articles/popularity/>

9/123

R's Roots... S

- S is a language that was developed by John Chambers and others at Bell Labs.
- S was initiated in 1976 as an internal statistical analysis environment - originally implemented as Fortran libraries.
- Early versions of the language did not contain functions for statistical modeling.
- In 1988 the system was rewritten in C and began to resemble the system that we have today (this was Version 3 of the language). The book *Statistical Models in S* by Chambers and Hastie (the blue book) documents the statistical analysis functionality.
- Version 4 of the S language was released in 1998 and is the version we use today. The book *Programming with Data* by John Chambers (the green book) documents this version of the language.

10/123

History of S

- In 1993 Bell Labs gave StatSci (now Insightful Corp.) an exclusive license to develop and sell the S language.
- In 2004 Insightful purchased the S language from Lucent for \$2 million and is the current owner.
- In 2006, Alcatel purchased Lucent Technologies and is now called Alcatel-Lucent.
- Insightful sells its implementation of the S language under the product name S-PLUS and has built a number of fancy features (GUIs, mostly) on top of it-hence the "PLUS".
- In 2008 Insightful is acquired by TIBCO for \$25 million; future of S-PLUS is uncertain.
- The S language itself has not changed dramatically since 1998.
- In 1998, S won the Association for Computing Machinery's Software System Award.

11/123

In "Stages in the Evolution of S", John Chambers writes:

"[W]e wanted users to be able to begin in an interactive environment, where they did not consciously think of themselves as programming. Then as their needs became clearer and their sophistication increased, they should be able to slide gradually into programming, when the language and system aspects would become more important."

<http://www.stat.bell-labs.com/S/history.html>

12/123

History of R

- 1991: Created in New Zealand by Ross Ihaka and Robert Gentleman. Their experience developing R is documented in a 1996 JCGS paper.
- 1993: First announcement of R to the public.
- 1995: Martin Mächler convinces Ross and Robert to use the GNU General Public License to make R free software.
- 1996: A public mailing list is created (R-help and R-devel)
- 1997: The R Core Group is formed (containing some people associated with S-PLUS). The core group controls the source code for R.
- 2000: R version 1.0.0 is released.
- 2013: R version 3.0 is released on April 3, 2013.
- Current version of R is 3.5.2
- As of February 14, 2019 `prettyNum(nrow(available.packages()), big.mark = ",")` packages listed on CRAN.

13/123

Euler's Formula

$$e^{i\pi} + 1 = 0$$

"The most remarkable formula in mathematics"
- Richard Feynman

16/123

Installing R

The latest version of R can be downloaded from cran.r-project.org. The current version of R is:

R.version\$version.string

```
## [1] "R version 3.5.2 (2018-12-20)"
```

You will also want to install [RStudio](https://www.rstudio.com/).

Installation instructions are available here:
<https://github.com/jbryer/introR/blob/master/Installation/Install.md>

To install the set of packages used for this workshop, run the following R command:

```
source("https://raw.githubusercontent.com/jbryer/introR/master/Installation/Setup.r")
```

14/123

R Packages

One aspect that makes R popular is how (relatively) easy it is to extend it's functionality vis-à-vis R packages. R packages are collections of R functions, data, and documentation.

The Comprehensive R Archive Network ([CRAN](https://cran.r-project.org/)) is the central repository where R packages are published. However, it should be noted that there are mirrors located across the globe.

Using packages requires two steps: first, install the package (required once per R installation); and second, load the package (once per R session).

```
install.packages("likert")
```

```
library(likert)
```

17/123

R as a Big Calculator

```
2 + 2
```

```
## [1] 4
```

```
1 + sin(9)
```

```
## [1] 1.412118
```

```
exp(i) ~ (1i + pi)
```

```
## [1] -1+0i
```

15/123

Installed and Loaded Packages

The `library()` function without any parameters will print all installed R packages whereas the `search()` function will list loaded packages (technically all available namespaces/environments, more on that later).

```
library()
search()
```

```
## [1] ".GlobalEnv"          "package:likert"      "package:stats"      "package:rbindbpy"
## [5] "package:scholar"       "package:lubridate"   "package:forcats"    "package:stringr"
## [9] "package:purrr"         "package:readr"      "package:tidyr"      "package:tidbits"
## [13] "package:tidyverse"     "package:dpbpyr"     "package:gplot2"     "package:stata"
## [17] "package:graphics"     "package:rtdrivers"  "package:utils"      "package:datasets"
## [21] "package:methods"      "AutoLoad"          "package:base"
```

18/123

Github Packages

Github is an online source repository and has become a popular place for R package developers to store their R packages. The `devtools` R package, designed to help package developers, has a function, `install_github` that will install R packages from a Github repository.

```
devtools::install_github("javier/likert")
```

19/123

ls()

We can use the `ls()` function to determine what functions are available in a package.

```
ls("package:likert")
```

```
## [1] "likert"          "likert.bar.plot"  "likert.density.plot" "likert.heat.plot"
## [5] "likert.histogram.plot" "likert.options"   "rcode"               "reverse.levels"
```

20/123

Getting Help

R provides extensive documentation and help. The `help.start()` function will launch a webpage with links to: * The R manuals * The R FAQ * Search engine * and many other useful sites

The `help.search()` function will search the help file for a particular word or phrase. For example:

```
help.search("cross tab")
```

To get documentation on a specific function, the `help()` function, or simply `?FunctionName` will open the documentation page in the web browser.

Lastly, to search the R mailing lists, use the `RSiteSearch()` function.

21/123

Reading Data

Data File Type	Extension	Function
R Data	rda, rdata	<code>base::load</code>
Comma separated values	csv	<code>utils::read.csv</code> , <code>readr::read_csv</code>
Other delimited files		<code>utils::read.table</code> , <code>readr::read_delim</code>
Tab separated files		<code>readr::read_tsv</code>
Fixed width files		<code>utils::read.fwf</code> , <code>readr::read_fwf</code>
SPSS	sav	<code>foreign::read.spss</code> , <code>haven::read_sav</code> , <code>haven::read_por</code>
SAS	sas	<code>haven::read_sas</code>
Read lines		<code>base::readLines</code> , <code>readr::read_lines</code>
Microsoft Excel	xls, xlsx	<code>xlsx::read.xlsx</code> , <code>readr::read_excel</code>
Qualtrics	h5p	<code>qualtrics::getSurveyResults</code>

23/123

Reading Data from Databases

The `RODBC` package is the most common way to connect to a variety of databases.

- `odbcConnect` - Open a connection to an ODBC database
- `sqlFetch` - Read a table from an ODBC database into a data frame
- `sqlQuery` - Submit a query to an ODBC database and return the results
- `close` - Close the connection

Other packages used to connect to specific databases:

- MySQL - `RMySQL`
- Oracle - `ROracle`
- JDBC - `RJDBC`
- SQLite - `RSQLite`
- PostgreSQL - `RPostgreSQL`

24/123

Loading Data

The `sqlutils` Package

The `sqlutils` is designed to help manage many query files and facilitates documenting and parameterizing the queries.

```
library(sqlutils)
sqlPath()

## [1] "/Users/joryer/Library/R/3.5/library/sqlutils/sql"

getQueries()

## [1] "StudentsInRange" "StudentSummary"

getParameters("StudentsInRange")

## [1] "startDate" "endDate"
```

25/123

Example SQL File (`StudentsInRange`)

```
#' Students enrolled within the given date range.
#'
#' @param startDate the start of the date range to return students.
#' @default startDate format(Sys.Date()), "%Y-01-01")
#' @param endDate the end of the date range to return students.
#' @default endDate format(Sys.Date()), "%Y-%m-%d")
#' Return CreatedDate the date the row was added to the warehouse data.
#' Return StudentId the student id.
SELECT *
FROM students
WHERE CreatedDate >= 'startDate' AND CreatedDate <= 'endDate';

sqlDoc("StudentsInRange")

#' Students enrolled within the given date range.
#' Parameters:
#'   param      desc      default
#'   startDate the start of the date range to return students.      '2012-01-01'
#'   endDate    the end of the date range to return students. format(Sys.Date()), "%Y-%m-%d")
#' default.val
#'   2012-01-01
#'   2019-02-14
#' Returns (note that this list may not be complete):
#'   variable      desc
#'   CreatedDate the date the row was added to the warehouse data.
#'   StudentId    the student id.
```

26/123

Running Query

```
require(RSQLite)
sqlFile <- paste(system.file(package="sqlutils"), "/db/students.db", sep="")
m <- dbDriver("SQLite")
conn <- dbConnect(m, dbname=sqlFile)
q1 <- sqlQuery("StudentSummary", connection=conn)
head(q1)

## CreatedDate count
## 1 2011-07-15 2886
## 2 2011-08-15 2983
## 3 2011-09-15 3071
## 4 2011-10-15 3059
## 5 2011-11-15 3058
## 6 2011-12-15 3074
```

27/123

IPEDS Data

The `ipeds` R package provides an interface to download data file from IPEDS.

```
library(ipeds)
data(surveys)
unique(surveys$Survey)

## [1] Institutional Characteristics      Enrollments
## [3] Completions                        Instructional staff/salaries
## [5] Full Staff                        Employees by Assigned Position
## [7] Finance                          Graduation Rates
## [9] Student Financial Aid and Net Price Admission and Test Score
## [10] Levels: Admission and Test Score Completions Employees by Assigned Position ... Student Financial Aid and Net

head(surveys[,c("SurveyID", "Title")])

## SurveyID      Title
## 1 HD            Directory Information
## 2 IC_ED          Educational offerings, organization, admissions, services and athletic associations
## 3 IC_AY          Student charges for academic year programs
## 4 IC_PV          Student charges by program (vocational programs)
## 5 FLAG          Response status for all survey components
## 6 SEST          Estimated enrollment
```

28/123

IPEDS Data

The `getIPEDSSurvey` and `ipedsHelp` are the most commonly used functions. The former will download and load the data into R (note data is cached and downloaded once per installation); the latter will provide the data dictionary for the given survey.

```
directory = getIPEDSSurvey("HS", 2013)
admissions = getIPEDSSurvey("IC", 2013)
retention = getIPEDSSurvey("ZPD", 2013)
```

ipedsHelp("HS", 2013)

head(directory)

```
## unitid      instnm      addr      city stabbr
## 1 100654      Alabama A & M University 4900 Meridian Street Normal AL
## 2 100663 University of Alabama at Birmingham Administration Bldg Suite 1070 Birmingham AL
## 3 100690      Armidge University      1200 Taylor Rd Montgomery AL
## 4 100704 University of Alabama in Huntsville 301 Sparkman Dr Huntsville AL
## 5 100724      Alabama State University 915 S Jackson Street Montgomery AL
## 6 100733 University of Alabama System Office 401 Queen City Ave Tuscaloosa AL
##   sip ftype obsorg      chfncm      chfctitle      gntsource      fctstale      ein
## 1 35762 1 5 Dr. Andrew Hugline, Jr. President 2.563725e+09 2163725030 636001109
## 2 35294-0110 1 5 Ray L. Watts President 2.059344e+09 2059377114 636005396
## 3 36117-3553 1 5 Michael Turner President 3.343877e+13 3343873878 237044224
## 4 35899 1 5 Robert A. Alankirch President 2.568246e+09 NA 636520830
## 5 31042-0271 1 5 Owen Boyd President 3.362229e+09 3363468661 636001103
## 6 35401 1 5 Robert Witt Chancellor 2.053490e+09 2053485206 636001138
##   opcid opcfay      website      admincl
## 1 100200 1 www.aamu.edu/ www.aamu.edu/admissions/pages/default.aspx
## 2 105200 1 www.uab.edu www.uab.edu/students/undergraduate-admissions
## 3 250400 1 www.amridgouniversity.edu www.amridgouniversity.edu/ua_admissions.html
## 4 105500 1 www.uab.edu admissions.uah.edu/
## 5 105500 1 www.alasu.edu/real/index.aspx www.alasu.edu/admissions/index.aspx
## 6 800400 2 www.uaSYSTEM.us.edu
##   faidurl
## 1 www.aamu.edu/admissions/finaidaid/Pages/default.aspx
## 2 www.uab.edu/students/paying-for-college
## 3 www.amridgouniversity.edu/ua_finaidaid.html
## 4 finaid.uah.edu/
## 5 www.alasu.edu/cont-aid/index.aspx/
## 6
## 1 www.aamu.edu/Adm/
## 2 uah.it.uab.edu/pla/etgprod/
## 3 https://www.amridgouniversity.edu/amridgop/login.aspx?ReturnUrl=/k2Emr/
## 4
## 5 psadmin.alasu.edu:8501/ppp/paged_1/EMPLOYEE/SHMS/c/ASU_SH_MENUD_MENU_ASU_SH_OHL_SECURE.GBL/PORTALPARAM_PCHNAV-3
## 6
## 6 upricurl sector
## 1 gallileo.aamu.edu/netpricemcalculator/epcalc.htm 1
## 2 www.collegeportraits.org/AL/UM/GetSector.aspx 1
## 3 tcc.noelllevits.com/(S(mwoihs4etsqpylf4dczuo))/Amridge420University/Freshman-Students 2
## 4 finaid.uah.edu/ 1
## 5 www.alasu.edu/cont-aid/forms/calculator/index.aspx/ 1
## 6 www.uaSYSTEM.us.edu 0
##   iclevel control hloffer ugoffer groffer hdegofr1 degyrant nbou hospital medical tribal locale
## 1 1 1 1 9 1 1 12 1 1 2 2 2 12
## 2 1 1 1 9 1 1 11 1 2 1 1 2 12
## 3 1 2 9 1 1 11 1 2 2 2 2 2 12
## 4 1 1 9 1 1 11 1 2 2 2 2 2 12
## 5 1 1 1 9 1 1 11 1 1 1 2 2 12
## 6 1 1 1 9 1 1 11 1 2 2 2 -2 2 13
```

29/123

The R Language

Arithmetic Operators

- `+` - addition
- `-` - subtraction
- `*` - multiplication
- `/` - division
- `^` or `**` - exponentiation
- `x % y` - modulus (x mod y) 5%%2 is 1

```
5 %% 2
## [1] 1

# x %/% y - integer division
5//2
## [1] 2
```

31/123

Lists

A **list** is an object that contains a list of named values

```
tmp <- list(a = 2112, b = pi, z = "Hello HEAIR!")
tmp

## $a
## [1] 2112
##
## $b
## [1] 3.141593
##
## $z
## [1] "Hello HEAIR!"

tmp[[1]] class(tmp[[1]]) # One square bracket: return a list

## $a
## [1] 2112

## [1] "list"

tmp[[1]] class(tmp[[2]]) # Two square brackets: return an object at that position

## [1] 2112

## [1] "numeric"
```

34/123

R Primitive Vectors

- **logical** (e.g. `TRUE`, `FALSE`)
- **integer** - whole numbers, either positive or negative (e.g. 2112, 42, -1)
- **double** or **numeric** - real number (e.g. 0.05, pi, -Inf, NaN)
- **complex** - complex number (e.g. 1i)
- **character** - sequence of characters, or a string (e.g. "Hello HEAIR!")

You can use the `class` function to determine the type of an object.

```
tmp <- c(2112, pi)
class(tmp)

## [1] "numeric"

To test if an object is of a particular class, use the is.XXX set of functions:

is.double(tmp)

## [1] TRUE

And to convert from one type to another, use the as.XXX set of functions:

as.character(tmp)

## [1] "2112"          "3.14159265358979"
```

32/123

Factors

A **factor** is a way for R to store a nominal, or categorical, variable. R stores the underlying data as an integer where each value corresponds to a label.

```
gender <- c(rep("male",4), rep("female", 6))
gender

## [1] "male" "male" "male" "male" "female" "female" "female" "female" "female" "female"

gender <- factor(gender, levels=c("male", "female", "unknown"))
gender


## [1] male male male male female female female female female female
## Levels: male female unknown

levels(gender)

## [1] "male" "female" "unknown"
```

35/123

Data Types in R



33/123

Factors can be ordered

The **ordered** parameter indicates whether the levels in the factor should be ordered.

```
library(TriMatch)

grade <- tutoring$Grade
table(grade, useMts="ifany")

## grade
## 5 1 2 3 4
## 187 25 86 271 573

## Loading required package: scales

## Attaching package: 'scales'

## The following object is masked from 'package:r':
## discard

## The following object is masked from 'package:readr':
## col_factor

## Loading required package: reshape2

## Attaching package: 'reshape2'

## The following object is masked from 'package:tidyr':
## smitha

## Loading required package: ez

data(tutoring, package="TriMatch")
head(tutoring$Grade)

## [1] 4 4 4 4 4 3

grade <- tutoring$Grade
table(grade, useMts="ifany")

## grade
## 5 1 2 3 4
## 187 25 86 271 573

grade <- factor(tutoring$Grade,
  levels=0:4,
  labels=c("F", "D", "C", "B", "A"),
  ordered=TRUE)
table(grade, useMts="ifany")

## grade
## F D C B A
## 187 25 86 271 573

With an ordered factor, coercing it back to an integer will maintain the order, but the values start with one!

head(grade)

## [1] A A A A A B
## Levels: F < D < C < B < A

table(as.integer(grade))

## 1 2 3 4 5
## 187 25 86 271 573

## [1] 4 4 4 4 4 3
```

36/123

Dates

R stores dates in **YYYY-MM-DD** format. The **as.Date** function will convert characters to **Dates** if they are in that form. If not, the **format** can be specified to help R coerce it to a **Date** format.

```
today <- Sys.Date()
format(today, "%a %d, %Y")
```

```
## [1] "February 14, 97"
```

```
as.Date("2015-MOV-01", format="%Y-%b-%d")
```

```
## [1] "2015-11-01"
```

- **%d** - day as a number (e.g. 0-31)
- **%a** - abbreviated weekday (e.g. **Mon**)
- **%A** - unabbreviated weekday (e.g. **Monday**)
- **%m** - month (i.e. 00-12)
- **%b** - abbreviated month (e.g. **Jan**)
- **%B** - unabbreviated month (e.g. **January**)
- **%y** - 2-digit year (e.g. 15)
- **%Y** - 4-digit year (e.g. 2015)

NA versus NULL

R is just as much a programming language as it is a statistical software package. As such it represents null differently for programming (using **NULL**) than for data (using **NA**).

- **NULL** represents the null object in R; it is a reserved word. **NULL** is often returned by expressions and functions whose values are undefined.
- **NA** is a logical constant of length 1 which contains a missing value indicator. **NA** can be freely coerced to any other vector type except **raw**. There are also constants **NA_integer**, **NA_real**, **NA_complex**, and **NA_character** of the other atomic vector types which support missing values: all of these are reserved words in the R language.

For more details, see <http://opendatagroup.com/2010/04/25/r-na-v-null/>

Handling Missing Data

There are a number of functions available for finding and subsetting missing values:

- **is.na** - function that takes one parameter and returns a logical vector of the same length where **TRUE** indicates the value is missing in the original vector.
- **complete.cases** - function that takes a data frame or matrix and returns **TRUE** if the entire row has no missing values.
- **na.omit** - function that takes a data frame and matrix and returns a subset of that data frame or matrix with any rows containing missing values removed.

Many statistical functions (e.g. **mean**, **sd**, **cor**) have a **na.rm** parameter that, when **TRUE**, will remove any missing values before calculating the statistic.

There are two very good R packages for imputing missing values:

- **mice** - [Multivariate Imputation by Chained Equations](#)
- **Amelia II** - [A Program for Missing Data](#)

37/123

Example: Excel File with Multiple Sheets

```
library(gdata)
summer2014 <- read.xls("MathMastery.xlsx", sheet=1)
fall2014 <- read.xls("MathMastery.xlsx", sheet=2)
summer2015 <- read.xls("MathMastery.xlsx", sheet=3)

summer2014$Term <- "Summer 2014"
fall2014$Term <- "Fall 2014"
summer2015$Term <- "Summer 2015"
```

```
mass <- rbind(summer2014, fall2014, summer2015)
head(mass)
```

```
##   Gender q1 q2 q3 q4 q5 q6 q7 q8 q9 q10 q11 q12 q13 q14      Term
## 1 Female 2 5 3 4 2 4 4 5 5 4 5 1 2 4 Summer 2014
## 2 Female 5 1 5 1 4 1 1 1 1 4 1 4 1 4 1 Summer 2014
## 3 Male 5 1 5 2 4 2 2 3 2 2 3 2 3 2 Summer 2014
## 4 Female 4 4 5 2 4 3 3 3 2 3 2 3 3 3 Summer 2014
## 5 Female 4 5 3 3 3 4 4 4 1 4 1 2 4 Summer 2014
## 6 Female 3 2 5 1 5 1 1 5 2 3 2 4 4 1 Summer 2014
```

Data Frames

Data frames are collection of vectors, thereby making them two dimensional. Unlike matrices (see **% matrix**) where all data elements are of the same type (i.e. numeric, character, logical, complex), each column in a data frame can be of a different type.

```
class(mass)
```

```
## [1] "data.frame"
```

```
dim(mass) # Dimension of the data frame (row by column)
```

```
## [1] 59 16
```

```
nrow(mass) # Number of rows
```

```
## [1] 59
```

```
ncol(mass) # Number of columns
```

```
## [1] 16
```

str

The **str** is perhaps the most useful function in R. It displays the structure of an R object.

```
str(mass)
```

```
## 'data.frame': 59 obs. of 16 variables:
## $ Gender: Factor w/ 2 levels "Female","Male": 1 1 2 1 1 1 1 1 1 ...
## $ q1 : int 2 5 5 4 4 5 4 4 5 3 ...
## $ q2 : int 5 1 1 4 5 2 2 4 4 5 ...
## $ q3 : int 3 5 5 5 5 5 5 5 3 ...
## $ q4 : int 4 1 2 2 3 1 1 3 2 5 ...
## $ q5 : int 2 4 4 4 3 5 5 4 2 2 ...
## $ q6 : int 4 1 2 3 3 1 2 5 4 5 ...
## $ q7 : int 4 1 2 3 4 1 1 4 4 5 ...
## $ q8 : int 5 1 3 3 4 5 3 5 4 5 ...
## $ q9 : int 5 1 2 2 4 2 1 4 4 5 ...
## $ q10 : int 4 4 2 3 1 3 2 3 3 1 ...
## $ q11 : int 5 1 2 4 2 1 4 4 5 ...
## $ q12 : int 1 4 3 3 1 4 4 2 3 1 ...
## $ q13 : int 2 4 3 3 2 4 4 3 3 1 ...
## $ q14 : int 4 1 2 3 4 1 2 4 2 5 ...
## $ Term : chr "Summer 2014" "Summer 2014" "Summer 2014" "Summer 2014" ...
```

40/123

41/123

42/123

Exploring the Data in Data Frames

```
head(mass)

##   Gender q1 q2 q3 q4 q5 q6 q7 q8 q9 q10 q11 q12 q13 q14      Term
## 1 Female 2 5 3 4 2 4 4 5 5 4 5 1 2 4 Summer 2014
## 2 Female 5 1 5 1 4 1 1 1 4 1 4 1 4 1 Summer 2014
## 3 Male 5 1 5 2 4 2 2 3 2 2 3 3 2 Summer 2014
## 4 Female 4 4 5 2 4 3 3 3 2 3 2 3 3 3 Summer 2014
## 5 Female 4 5 5 3 3 4 4 4 1 4 1 2 4 Summer 2014
## 6 Female 5 2 5 1 5 1 1 5 2 3 2 4 4 1 Summer 2014
```

```
tail(mass, n=3)
```

```
##   Gender q1 q2 q3 q4 q5 q6 q7 q8 q9 q10 q11 q12 q13 q14      Term
## 57 Female 1 5 1 5 1 2 5 5 5 4 4 1 1 5 Summer 2015
## 58 Male 4 3 5 2 5 2 2 3 2 5 2 3 4 3 Summer 2015
## 59 Male 5 1 5 1 5 1 1 3 1 5 1 5 5 1 Summer 2015
```

The **View** function will provide a (read-only) spreadsheet view of the data frame.

```
View(mass)
```

43/123

Tip: One Column Data Frame

When selecting one column from a data frame, R will convert the returned object to a vector.

```
class(mass[,1])
```

```
## [1] "factor"
```

You can use the **drop=FALSE** parameter keep the subset as a data frame.

```
class(mass[,1,drop=FALSE])
```

```
## [1] "data.frame"
```

46/123

Subsetting Data Frames

Using square brackets will allow you to subset from a data frame. The first parameter is for rows, the second for columns. Leaving one blank will return all rows or columns.

```
mass[c(1:2,10),] # Return the first, second, and tenth row
```

```
##   Gender q1 q2 q3 q4 q5 q6 q7 q8 q9 q10 q11 q12 q13 q14      Term
## 1 Female 2 5 3 4 2 4 4 5 5 4 5 5 2 4 Summer 2014
## 2 Female 5 1 5 1 4 1 1 1 4 1 4 4 1 Summer 2014
## 10 Female 1 5 3 5 2 5 5 5 5 1 5 1 1 5 Summer 2014
```

```
mass[,2] # Return the second column
```

```
## [1] 2 5 5 4 4 5 4 5 1 3 4 4 1 5 2 2 4 3 5 1 2 3 2 4 4 4 1 4 2 5 4 4 3 3 4 5 3 3 1 4 4 3 1 4 4
## [49] 5 5 3 4 3 2 4 5 1 4 5
```

You can also subset columns using the dollar sign (\$) notation.

```
mass$q10
```

```
## [1] 4 4 2 3 1 2 3 3 1 1 2 4 1 5 2 2 3 3 4 4 3 2 3 3 5 4 4 1 3 3 5 4 2 3 2 5 5 2 3 1 3 4 1 4 4
## [49] 5 3 1 4 1 3 4 4 4 5 5
```

44/123

Subsetting with Logical Operators

You can subset using logical vectors. For example, there are 7764 rows in the **directory** data frame loaded from IPEDS. You can pass a logical vector of length 7764 where **TRUE** indicates to return the row and **FALSE** to not. For example, we wish to return the row with Excelsior College:

```
row <- directory$instnm == "Excelsior College"
length(row)
```

```
## [1] 7764
```

Here we are using the == logical operator. This will test each element in the **directory\$instnm** and return **TRUE** if it is equal to **Excelsior College**, **FALSE** otherwise.

```
directory[row, 1:16] # Include only 16 columns for display purposes
```

```
##   unitid      instnm      addr      city stabbr      zip fips obereg      chfcm
## 2783 196680 Excelsior College 7 Columbia Cir Albany NY 12203-5159 36 2 John P. Ebersole
## chftitle gentele fastele min opaid opelag webaddr
## 2783 President 5184448500 5184448777 141599643 203400 1 www.excelsior.edu
```

47/123

Subsetting Missing Values

Using the **complete.cases** function, we can return rows with at least one missing values.

```
mass[complete.cases(mass),]
```

```
##   Gender q1 q2 q3 q4 q5 q6 q7 q8 q9 q10 q11 q12 q13 q14      Term
## 26 Female 3 4 NA 3 2 4 2 4 3 2 4 2 2 4 Fall 2014
## 38 Male 4 2 5 2 5 2 2 3 2 5 NA 3 4 3 Fall 2014
## 53 Female 3 3 NA 3 2 3 3 4 4 1 3 2 3 4 Summer 2015
```

Using the **is.na**, we can change replace the missing values.

```
(tmp <- sample(c(1:5, NA)))
```

```
## [1] 5 3 2 1 NA 4
```

```
tmp[is.na(tmp)] <- 2112
tmp
```

```
## [1] 5 3 2 1 2112 4
```

45/123

which

The **which** command will return an **integer** vector with the positions within the **logical** vector that are **TRUE**.

```
which(row)
```

```
## [1] 2783
```

```
directory[2783, 1:16]
```

```
##   unitid      instnm      addr      city stabbr      zip fips obereg      chfcm
## 2783 196680 Excelsior College 7 Columbia Cir Albany NY 12203-5159 36 2 John P. Ebersole
## chftitle gentele fastele min opaid opelag webaddr
## 2783 President 5184448500 5184448777 141599643 203400 1 www.excelsior.edu
```

48/123

Logical Operators

- **!a** - TRUE if a is FALSE
- **a == b** - TRUE if a and b are equal
- **a != b** - TRUE if a and b are not equal
- **a > b** - TRUE if a is larger than b, but not equal
- **a >= b** - TRUE if a is larger or equal to b
- **a < b** - TRUE if a is smaller than b, but not equal
- **a <= b** - TRUE if a is smaller or equal to b
- **a %in% b** - TRUE if a is in b where b is a vector

```
which( letters %in% c("a","e","i","o","u") )
```

```
## [1] 1 5 9 15 21
```

- **a | b** - TRUE if a or b are TRUE
- **a & b** - TRUE if a and b are TRUE
- **isTRUE(a)** - TRUE if a is TRUE

49/123

Reshaping Data (melting)

Data is often said to be in one of two formats: wide or long. The **mass** data frame is currently in a wide format where each variable is a separate column. However, there are certain analyses that will require the data to be in a long format. In a long format, we would have two columns to represent all the items (one for the item name, one for value), plus any additional identity variables. The **melt** command will convert a wide table to a long table.

```
library(reshape2)
mass$Id <- 1:nrow(mass) # 59 rows
mass.melted <- melt(mass, id.vars=c("Id","Gender","Term"), variable.name="Item", value.name="Response")
head(mass.melted, n=4)

##   Id Gender   Term Item Response
## 1 1 Female Summer 2014 q1      2
## 2 2 Female Summer 2014 q1      5
## 3 3 Male Summer 2014 q1      5
## 4 4 Female Summer 2014 q1      4
```

```
nrow(mass.melted)
```

```
## [1] 826
```

52/123

Side Note: Operators are Functions

All operations (e.g. +, -, *, /, [, <-) are functions.

```
class("+")
```

```
## [1] "function"
```

```
"+"
```

```
## function (e1, e2) .Primitive("+")
```

```
"+"(2, 3)
```

```
## [1] 5
```

You can redefine these functions, but probably not a good idea :-)

50/123

Reshaping Data (casting)

To convert a long table to a wide table, use the **dcast** function

```
mass.casted <- dcast(mass.melted, Id + Gender + Term ~ Item, value.var="Response")
head(mass.casted)
```

```
##   Id Gender   Term q1 q2 q3 q4 q5 q6 q7 q8 q9 q10 q11 q12 q13 q14
## 1 1 Female Summer 2014 2 5 3 4 2 4 4 5 5 4 5 1 2 4
## 2 2 Female Summer 2014 5 1 5 1 4 1 1 1 1 4 1 4 4 1
## 3 3 Male Summer 2014 5 1 5 2 4 2 2 3 2 2 2 3 3 2
## 4 4 Female Summer 2014 4 4 5 2 4 3 3 3 2 3 2 3 3 3
## 5 5 Female Summer 2014 4 5 5 3 3 3 4 4 4 1 4 1 2 4
## 6 6 Female Summer 2014 5 2 5 1 5 1 1 5 2 3 2 4 4 1
```

```
## [1] 59
```

53/123

Sorting Data

The **order** function will take one or more vectors (usually in the form of a data frame) and return an integer vector indicating the new order. There are two parameters to adjust where **NA**s are placed (**na.last=FALSE**) and whether to sort in increasing or decreasing order (**decreasing=FALSE**).

```
(randomLetters <- sample(letters))
```

```
## [1] "q" "l" "e" "h" "q" "e" "a" "a" "a" "a" "a" "g" "p" "m" "b" "j" "l" "d" "u" "u" "e" "a" "a" "a"
## [25] "a" "t"
```

```
randomLetters[order(randomLetters)]
```

```
## [1] "a" "a" "a" "d" "u" "e" "g" "h" "l" "j" "a" "l" "a" "a" "a" "p" "q" "p" "a" "a" "u" "u" "a" "a"
## [25] "p" "a"
```

```
randomLetters[order(randomLetters, decreasing=TRUE)]
```

```
## [1] "q" "p" "a" "u" "u" "a" "a" "a" "g" "p" "a" "a" "l" "a" "j" "l" "a" "g" "e" "a" "d" "a"
## [25] "b" "a"
```

51/123

Removing Columns

To remove a single column from a data frame, simply assign to **NULL** to the column value.

```
mass$Id <- NULL
head(mass)
```

```
##   Gender q1 q2 q3 q4 q5 q6 q7 q8 q9 q10 q11 q12 q13 q14   Term
## 1 Female 2 5 3 4 2 4 4 5 5 4 5 1 2 4 Summer 2014
## 2 Female 5 1 5 1 4 1 1 1 1 4 1 4 1 Summer 2014
## 3 Male 5 1 5 2 4 2 2 3 2 2 2 3 2 Summer 2014
## 4 Female 4 4 5 2 4 3 3 3 2 3 2 3 3 Summer 2014
## 5 Female 4 5 5 3 3 3 4 4 4 1 4 1 2 4 Summer 2014
## 6 Female 5 2 5 1 5 1 1 5 2 3 2 4 4 1 Summer 2014
```

54/123

if Statements

If statements are a way of doing different operations conditionally. In this example, the **sample** function returns either a 0 or 1.

```
if(sample(c(0,1), 1) > 0.5) {  
  print("heads")  
} else {  
  print("tails")  
}
```

```
## [1] "heads"
```

For simpler problems, the **ifelse** function may be easier.

```
ifelse(sample(c(0,1), 1) > 0.5, 'heads', 'tails')
```

```
## [1] "tails"
```

55/123

Functions

Functions allow for organizing common procedures to easily be used later. You can specify any number of parameters with optional default values. The objects created and/or edited within the function are local to the function (i.e. not available to the calling environment). The **return** or **invisible** function is used to return value(s) to the environment calling the function. The **invisible** function will return a value but will not print it to the console if it is not assigned to an object.

```
getStanding <- function(credits,  
  breaks=c(30, 60, 90),  
  labels=c("Freshman", "Sophomore", "Junior", "Senior")) {  
  results <- cut(credits, breaks=c(-Inf, breaks, Inf), labels=labels,  
    include.lowest=TRUE, right=FALSE)  
  return(results)  
}
```

```
getStanding(42)
```

```
## [1] Sophomore  
## Levels: Freshman Sophomore Junior Senior
```

```
getStanding(c(27, 60, 101, 200))
```

```
## [1] Freshman Junior Senior Senior  
## Levels: Freshman Sophomore Junior Senior
```

58/123

for loops

The **for** loop allows you to do some operation a number of times. In this example, we wish to convert each of the Likert responses in the **mass** data frame to an ordered factor.

```
for(i in 2:15) {  
  mass[,i] <- factor(mass[,i], levels=c(1,5, labels=c("Strongly Disagree",  
    "Disagree", "Neutral", "Agree", "Strongly Agree"), ordered=TRUE)  
}  
str(mass)
```

```
## 'data.frame': 59 obs. of 16 variables:  
## $ Gender: Factor w/ 2 levels "Female","Male": 1 1 2 1 1 1 1 1 1 1 ...  
## $ q1 : Ord.factor w/ 5 levels "Strongly Disagree"<... 2 5 5 4 4 5 4 5 1 ...  
## $ q2 : Ord.factor w/ 5 levels "Strongly Disagree"<... 5 1 4 5 2 2 4 4 5 ...  
## $ q3 : Ord.factor w/ 5 levels "Strongly Disagree"<... 3 5 5 5 5 5 5 5 3 ...  
## $ q4 : Ord.factor w/ 5 levels "Strongly Disagree"<... 4 1 2 2 1 1 3 2 5 ...  
## $ q5 : Ord.factor w/ 5 levels "Strongly Disagree"<... 2 4 4 3 5 5 4 2 2 ...  
## $ q6 : Ord.factor w/ 5 levels "Strongly Disagree"<... 4 1 2 3 1 3 2 5 4 ...  
## $ q7 : Ord.factor w/ 5 levels "Strongly Disagree"<... 4 1 2 3 4 1 4 4 5 ...  
## $ q8 : Ord.factor w/ 5 levels "Strongly Disagree"<... 5 1 3 4 5 3 5 4 5 ...  
## $ q9 : Ord.factor w/ 5 levels "Strongly Disagree"<... 5 1 2 4 2 1 4 4 5 ...  
## $ q10 : Ord.factor w/ 5 levels "Strongly Disagree"<... 4 4 2 3 1 3 2 3 1 ...  
## $ q11 : Ord.factor w/ 5 levels "Strongly Disagree"<... 5 1 2 4 3 1 4 4 5 ...  
## $ q12 : Ord.factor w/ 5 levels "Strongly Disagree"<... 1 4 3 3 1 4 4 2 3 1 ...  
## $ q13 : Ord.factor w/ 5 levels "Strongly Disagree"<... 2 4 3 2 4 4 3 3 1 ...  
## $ q14 : Ord.factor w/ 5 levels "Strongly Disagree"<... 4 1 2 3 4 1 2 4 2 5 ...  
## $ Term : chr "Summer 2014" "Summer 2014" "Summer 2014" "Summer 2014" ...
```

56/123

Renaming Columns

```
names(mass) # Get the current names
```

```
## [1] "Gender" "q1" "q2" "q3" "q4" "q5" "q6" "q7" "q8" "q9"  
## [11] "q10" "q11" "q12" "q13" "q14" "Term"
```

```
items <- c('I find math interesting.',  
  'I get uptight during math tests.',  
  'I think that I will use math in the future.',  
  'Mind goes blank and I am unable to think clearly when doing my math test.',  
  'Math relates to my life.',  
  'I worry about my ability to solve math problems.',  
  'I get a sinking feeling when I try to do math problems.',  
  'I find math challenging.',  
  'Mathematics makes me feel nervous.',  
  'I would like to take more math classes.',  
  'Mathematics makes me feel uneasy.',  
  'Math is one of my favorite subjects.',  
  'I enjoy learning with mathematics.',  
  'Mathematics makes me feel confused.')
```

```
names(mass) <- c("Gender", items, "Term")
```

59/123

while loops

while loops will run until some condition is met. In this example, consider 1=heads, 0=tails. How many random events would it take before getting 100 heads?

```
sum <- 0  
count <- 0  
while(sum < 100) {  
  sum <- sum + sample(c(0,1), 1)  
  count <- count + 1  
}  
count # Number of loops until we got 100 head
```

```
## [1] 214
```

57/123

Example: SAT and First Year Retention (subsetting)

In this example, we wish to explore the relationship between SAT scores and first year retention as measures at the institutional level. These data are part of the IPEDS data collection, but are collected in different surveys. The first step is to subset the data frames so we are working with fewer columns. This is not necessary, but simplifies the analysis.

```
directory <- directory[c("unitid", "instnm", "sector", "control")]  
retention <- retention[,c("unitid", "ret_pct", "ret_pct")]  
admissions <- admissions[,c("unitid", "admcst", "admcst", "admcst", "admcst",  
  "appcsm", "appcsm", "admcst", "admcst", "admcst",  
  "enrlftw", "enrlftw", "enrlftw", "enrlftw", "enrlftw",  
  "actmat", "actmat", "actmat", "actmat", "actmat",  
  "actwr75", "actwr75", "actwr75", "actwr75", "actwr75",  
  "actact25", "actact25", "actact25", "actact25", "actact25",  
  "actwr75", "actwr75", "actwr75", "actwr75", "actwr75")]
```

60/123

Example: SAT and First Year Retention (converting factors)

Next, we will recode the variables that indicate whether SAT scores are required for admission.

```
admissionsLabels <- c("Required", "Recommended", "Neither required nor recommended",  
  "Do not know", "Not reported", "Not applicable")  
admissions$admcn1 <- factor(admissions$admcn1, levels=c(1,2,3,4,-1,-2),  
  labels=admissionsLabels)  
admissions$admcn2 <- factor(admissions$admcn2, levels=c(1,2,3,4,-1,-2),  
  labels=admissionsLabels)  
admissions$admcn7 <- factor(admissions$admcn7, levels=c(1,2,3,4,-1,-2),  
  labels=admissionsLabels)
```

61/123

Example: SAT and First Year Retention (converting numeric columns)

IPEDS uses periods (.) to represent missing values. As a result, R will treat the column as a **character** column so we need to convert them to **numeric** columns. The **as.numeric** function will do this and any value that is not numeric (.s in this example) will be treated as missing values (i.e. **NA**).

```
ret2$SATmath75 <- as.numeric(ret2$SATmath75)  
ret2$SATmath25 <- as.numeric(ret2$SATmath25)  
ret2$SATwriting75 <- as.numeric(ret2$SATwriting75)  
ret2$SATwriting25 <- as.numeric(ret2$SATwriting25)  
ret2$HunSATScores <- as.integer(ret2$HunSATScores)
```

64/123

Example: SAT and First Year Retention (renaming variables)

Next, rename the variables to more understandable names.

```
names(retention) <- c("unitid", "FullTimeRetentionRate", "PartTimeRetentionRate")  
names(admissions) <- c("unitid", "useEDPA", "useEDPA", "useAdmissionTestScore",  
  "ApplicantsMen", "ApplicantsWomen", "ApplicantsTotal",  
  "AdmissionsMen", "AdmissionsWomen", "AdmissionsTotal",  
  "EnrolledFullTimeMen", "EnrolledFullTimeWomen",  
  "EnrolledPartTimeMen", "EnrolledPartTimeWomen",  
  "EnrolledTotal", "HunSATScore", "PercentSATScore",  
  "HunACTScore", "PercentACTScore", "SATReading25",  
  "SATReading75", "SATMath25", "SATMath75", "SATWriting25",  
  "SATWriting75", "ACTComposite25", "ACTComposite75",  
  "ACTEngliah25", "ACTEngliah75", "ACTMath25", "ACTMath75",  
  "ACTWriting25", "ACTWriting75")
```

62/123

Example: SAT and First Year Retention (calculating)

IPEDS only provides the 25th and 75th percentile in SAT and ACT scores. We will use the mean of these two values as a proxy for the mean.

```
ret2$SATmath <- (ret2$SATmath75 + ret2$SATmath25) / 2  
ret2$SATwriting <- (ret2$SATwriting75 + ret2$SATwriting25) / 2  
ret2$SATtotal <- ret2$SATmath + ret2$SATwriting  
ret2$acceptanceTotal <- as.numeric(ret2$AdmissionsTotal) / as.numeric(ret2$ApplicantsTotal)  
ret2$useAdmissionTestScore <- as.factor(as.character(ret2$useAdmissionTestScore))
```

65/123

Example: SAT and First Year Retention (merging)

We need to merge the three data frames to a single data frame. The **merge** function will merge, or join, two data frames on one or more columns. In this example schools that do not appear in all three data will not appear in the final data frame. To control how data frames are merge, see the **all**, **all.x**, and **all.y** parameters of the **merge** function (hint: works like outer joins in SQL).

```
ret <- merge(directory, admissions, by="unitid")  
ret <- merge(ret, retention, by="unitid")
```

We will also only use schools that require or recommend admission tests.

```
ret2 <- ret[ret$useAdmissionTestScore !=  
  c("Required", "Recommended", "Neither required nor recommended"),]
```

63/123

Example: SAT and First Year Retention (final data frame)

```
str(ret2)  
  
## 'data.frame': 2281 obs. of 42 variables:  
## $ unitid : int 100654 100663 100706 100724 100751 100830 100858 100937 101116 101365 ...  
## $ instnm : chr "Alabama A & M University" "University of Alabama at Birmingham" "University of J  
## $ sector : int 1 1 1 1 1 1 1 2 3 3 ...  
## $ control : int 1 1 1 1 1 1 1 2 3 3 ...  
## $ useEDPA : Factor w/ 6 levels "Required", "Recommended",...: 2 3 2 3 2 2 1 3 3 ...  
## $ useSRnak : Factor w/ 3 levels "Required", "Recommended",...: 3 3 3 3 3 2 3 2 2 ...  
## $ useAdmissionTestScore : chr "2401" "2214" "1183" "3808" ...  
## $ ApplicantsMen : chr "2741" "3475" "871" "6436" ...  
## $ ApplicantsWomen : chr "1442" "5689" "2054" "1245" ...  
## $ ApplicantsTotal : chr "2108" "1944" "978" "1813" ...  
## $ AdmissionsMen : chr "3421" "2990" "661" "3438" ...  
## $ AdmissionsWomen : chr "3521" "4934" "1456" "5215" ...  
## $ AdmissionsTotal : chr "533" "4897" "387" "196" ...  
## $ EnrolledFullTimeMen : chr "556" "1835" "251" "840" ...  
## $ EnrolledFullTimeWomen : chr "9" "20" "9" "24" ...  
## $ EnrolledPartTimeMen : chr "6" "21" "4" "19" ...  
## $ EnrolledPartTimeWomen : chr "1184" "1773" "631" "1479" ...  
## $ EnrolledTotal : int 167 103 118 269 1469 0 617 113 NA NA ...  
## $ PercentSATScore : chr "15" "6" "34" "18" ...  
## $ PercentACTScore : chr "968" "1645" "610" "1285" ...  
## $ PercentSATScore : chr "88" "93" "94" "97" ...  
## $ SATReading25 : chr "370" "520" "510" "380" ...  
## $ SATReading75 : chr "450" "640" "640" "480" ...  
## $ SATmath25 : num 350 320 510 370 500 NA 540 520 NA NA ...  
## $ SATmath75 : num 450 650 650 480 640 NA 650 640 NA NA ...  
## $ SATwriting25 : num NA NA NA NA 480 NA 510 NA NA ...  
## $ SATwriting75 : num NA NA NA NA 600 NA 620 NA NA ...  
## $ ACTComposite25 : chr "15" "22" "23" "15" ...  
## $ ACTComposite75 : chr "19" "28" "29" "19" ...  
## $ ACTEngliah25 : chr "14" "22" "22" "14" ...  
## $ ACTEngliah75 : chr "19" "29" "30" "20" ...  
## $ ACTmath25 : chr "15" "20" "22" "15" ...  
## $ ACTmath75 : chr "18" "24" "28" "18" ...  
## $ ACTwriting25 : chr "." "." "." ...  
## $ ACTwriting75 : chr "." "." "." ...  
## $ FullTimeRetentionRate : int 43 80 81 62 87 43 89 80 19 50 ...  
## $ PartTimeRetentionRate : int 50 50 44 39 46 43 81 NA 14 NA ...  
## $ SATmath : num 400 585 580 425 570 NA 595 580 NA NA ...  
## $ SATwriting : num NA NA NA NA 540 NA 540 NA 545 NA NA ...  
## $ SATtotal : num NA NA NA NA 1110 NA 1160 NA NA ...  
## $ AcceptanceTotal : num 0.899 0.867 0.806 0.513 0.565 ...
```

66/123

Miscellaneous Functions

• **paste** and **paste0** - concatenate strings (**paste0** uses **sep=""** by default)

```
paste("Hello", "HEALTH")
```

```
## [1] "Hello HEALTH"
```

• **prettyNum** - Formats numbers to strings

```
prettyNum(123456.987654321, big.mark="," , digits=8)
```

```
## [1] "123,456.99"
```

67/123

The **describe** and **describeBy** Functions

The **describe** and **describeBy** functions in the **psych** package are useful for calculating descriptive statistics across many variables.

```
library(psych)
describe(tutoringAge)
```

```
##      vars      n mean      sd median trimmed      mad min max range      skew      kurtosis      se
## X1      1 1142 36.92 9.05      37  36.51 10.38  20  65      45 0.36      -0.49 0.27
```

```
describeBy(tutoringAge, tutoringMilitary, mat=TRUE)
```

```
##      item group1      vars      n      mean      sd median      trimmed      mad min max range      skew      kurtosis
## X11      1 FALSE      1 783 38.42529 9.472897      38 38.15470 10.3782 20 65      45 0.2089516
## X12      2  TRUE      1 359 33.64903 7.015899      33 33.40484  7.4130 21 53      32 0.2938581
##      kurtosis      se
## X11 -0.6754505 0.3385337
## X12 -0.5781302 0.3702850
```

The **mat=TRUE** parameter presents the results in tabular format that is useful for later converting to a document or report.

70/123

Statistics

71/123

Contingency and Proportion Tables

```
(tab1 <- table(tutoringMilitary, useRna="ifany"))
```

```
##
##      Control  Treat1  Treat2
## FALSE TRUE
## 783 359
```

```
(tab2 <- table(tutoringMilitary, tutoringTreat,
               useRna="ifany"))
```

```
##
##      Control  Treat1  Treat2
## FALSE 609 102 72
## TRUE 309 32 18
```

```
prop.table(tab1)
```

```
##
##      Control  Treat1  Treat2
## FALSE TRUE
## 0.6856392 0.3143608
```

```
prop.table(tab2)      # Sum of all cells = 1
```

```
##
##      Control  Treat1  Treat2
## FALSE 0.53327496 0.08931699 0.06304729
## TRUE 0.27057793 0.02802102 0.01576182
```

```
prop.table(tab2, 1) # Sum of rows = 1
```

```
##
##      Control  Treat1  Treat2
## FALSE 0.77777778 0.13204820 0.09195402
## TRUE 0.86072423 0.08913649 0.05013928
```

```
prop.table(tab2, 2) # Sum of columns = 1
```

```
##
##      Control  Treat1  Treat2
## FALSE 0.6833987 0.7611940 0.8000000
## TRUE 0.3166013 0.2388060 0.2000000
```

71/123

Descriptive Statistics

```
mean(tutoringAge)      # Mean
```

```
## [1] 36.92282
```

```
median(tutoringAge)      # Median
```

```
## [1] 37
```

```
var(tutoringAge)      # Variance
```

```
## [1] 81.86711
```

```
sd(tutoringAge)      # Standard deviation
```

```
## [1] 9.048045
```

```
IQR(tutoringAge)      # Inter-quartile range
```

```
## [1] 13
```

```
min(tutoringAge)      # Minimum value
```

```
## [1] 20
```

```
max(tutoringAge)      # Maximum value
```

```
## [1] 65
```

```
summary(tutoringAge) # All at once
```

```
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##      20.00   30.00   37.00   36.92   43.00   65.00
```

Correlation

```
cor(ret2SATMath, ret2FullTimeRetentionRate, use="complete.obs")
```

```
## [1] 0.7320969
```

```
cor.test(ret2SATMath, ret2FullTimeRetentionRate, use="complete.obs")
```

```
##
## Pearson's product-moment correlation
##
## data: ret2SATMath and ret2FullTimeRetentionRate
## t = 38.987, df = 1316, p-value < 2.2e-16
## alternative hypothesis: true correlation is not equal to 0
## 95 percent confidence interval:
##  0.7060894 0.7562401
## sample estimates:
##      cor
## 0.7320969
```

72/123

t-tests

```
t.test(Grade ~ treat2, data=tutoring)

##
## Welch Two Sample t-test
##
## data: Grade by treat2
## t = -6.3194, df = 550.26, p-value = 5.427e-10
## alternative hypothesis: true difference in means is not equal to 0
## 95 percent confidence interval:
##  -0.6720923 -0.3533502
## sample estimates:
## mean in group FALSE mean in group TRUE
##           2.790850           3.303571
```

Regression

The `lm` and `glm` functions are used for linear models and generalized linear models, respectively. The following example is a multiple regression model predicting `FullTimeRetentionRate` from `SATWriting`, `SATMath`, `AcceptanceTotal`, and `UseAdmissionTestScores`.


```
lm.fit = lm(FullTimeRetentionRate ~ SATWriting + SATMath + AcceptanceTotal +
            UseAdmissionTestScores, data=ret2, weights=ret2$NumSATCourses)
summary(lm.fit)

##
## Call:
## lm(formula = FullTimeRetentionRate ~ SATWriting + SATMath + AcceptanceTotal +
##     UseAdmissionTestScores, data = ret2, weights = ret2$NumSATCourses)
##
## Weighted Residuals:
##      Min       1Q   Median       3Q      Max
## -904.06  -72.28  -16.70   42.73  632.65
##
## Coefficients:
##              (Intercept)              13.127102              3.499094              3.752
## SATWriting              0.037229              0.010084              3.692
## SATMath              0.005770              0.009463              9.104
## AcceptanceTotal              4.947624              1.394498              3.540
## UseAdmissionTestScoresRecommended              -4.236870              2.171028              -1.866
## UseAdmissionTestScoresRequired              -2.394511              2.116747              -1.121
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 133.1 on 790 degrees of freedom
## (1485 observations deleted due to missingness)
## Multiple R-squared:  0.7461, Adjusted R-squared:  0.7445
## F-statistic: 464.2 on 5 and 790 DF, p-value: < 2.2e-16
```


See [Linear Regression with NYS Report Card](#) for more linear regression.

Data Visualization


Histograms



Boxplots




Barplots



Scatter Plots

```
plot(rel2$SATMath, rel2$FullTimeBoutRate)
```



79/123


Gammar of Graphics with **ggplot2**

- **ggplot2** is an R package that provides an alternative framework based upon Wilkinson's (2005) Grammar of Graphics.
- **ggplot2** is, in general, more flexible for creating "prettier" and complex plots.
- Works by creating layers of different types of objects/geometries (i.e. bars, points, lines, polygons, etc.)
- **ggplot2** has at least three ways of creating plots:
 1. `qplot`
 2. `ggplot(...) + geom_XXX(...) + ...`
 3. `ggplot(...) + layer(...)`
- We will focus only on the second.

80/123

First Example

```
data(diamonds)
ggplot(diamonds, aes(x=carat, y=price, color=cut)) + geom_point()
```



81/123

Parts of a **ggplot2** Statement

- Data
`ggplot(myDataFrame, aes(x=x, y=y))`
- Layers
`geom_point()`, `geom_histogram()`
- Facets
`facet_wrap(~ cut)`, `facet_grid(~ cut)`
- Scales
`scale_y_log10()`
- Other options
`ggtitle('my title')`, `ylim(c(0, 10000))`, `xlab('x-axis label')`

82/123

Lots of geoms


```
ls('package:ggplot2')[grep('geom_', ls('package:ggplot2'))]
```

## [1] "geom_abline"	"geom_area"	"geom_bar"	"geom_bin2d"
## [5] "geom_blank"	"geom_boxplot"	"geom_coi"	"geom_contour"
## [9] "geom_count"	"geom_crossbar"	"geom_curve"	"geom_density"
## [13] "geom_density_2d"	"geom_density2d"	"geom_dotplot"	"geom_errorbar"
## [17] "geom_errorbarh"	"geom_frugpoly"	"geom_hex"	"geom_histogram"
## [21] "geom_jitter"	"geom_label"	"geom_line"	"geom_line2"
## [25] "geom_linerange"	"geom_map"	"geom_path"	"geom_point"
## [29] "geom_pointrange"	"geom_polygon"	"geom_qq"	"geom_qq_line"
## [33] "geom_quantile"	"geom_raster"	"geom_ribbon"	"geom_ridge"
## [37] "geom_rug"	"geom_segment"	"geom_sf"	"geom_sf_label"
## [41] "geom_sf_text"	"geom_smooth"	"geom_spoke"	"geom_step"
## [45] "geom_text"	"geom_tile"	"geom_violin"	"geom_vline"
## [49] "update_geom_defaults"			

83/123


Histograms

```
ggplot(tutoring, aes(x=Age)) + geom_histogram(binwidth=5)
```




84/123

Density Plots




85/123

Boxplots



86/123

Barplots



87/123

Example: SAT and First Year Retention

First, we define the data and aesthetics (i.e. the various components for the graph). Here, SAT math scores will be on the x-axis, full-time retention rates on the y-axis. Points will be sized based upon the number of SAT scores included. The color will represent whether admission test scores are required, recommended, or neither required nor recommended.

```
p <- ggplot(ret2, aes(x=SATMath, y=FullTimeRetentionRate, size=NumSATScores, color=UsedAdmissionTestScores))
```

Next, we define the graphic objects (i.e. `geoms`) to add to the figure.

```
p <- p + geom_point()
```

We can override the default color scheme for the points (note the default is `scale_color_hue`).


```
p <- p + scale_color_brewer('Test Required for Admissions', type='qual')
```

Lastly, we will add axis labels and a title.

```
p <- p + xlab('SAT Math Score') + ylab('Full Time Retention Rate') + ggtitle('SAT and Full Time Retention')
```

88/123

Example: SAT and First Year Retention




89/123

Example: SAT and First Year Retention

We can easily add additional geoms, here adding a Loess regression line.

```
p + geom_smooth(method='loess')
```



90/123


Likert Type Items

```
library(likert)
l <- Likert(masa[,2:15])
summary(l)

##              Item low neutral high mean
## 3      I think that I will use math in the future.    0    20    80 4.35
## 8      I find math challenging.                    15    15    70 3.75
## 1      I find math interesting.                     25    10    65 3.60
## 5      Math relates to my life.                     25    10    65 3.60
## 2      I get uptight during math tests.             30    20    50 3.35
## 14     Mathematics makes me feel confused.          35    15    50 3.15
## 6      I worry about my ability to solve math problems. 40    20    40 2.90
## 7      I get a sinking feeling when I try to do math problems. 45    15    40 2.65
## 9      Mathematics makes me feel nervous.           45    15    40 2.90
## 11     Mathematics makes me feel uneasy.            40    30    30 2.85
## 13     I enjoy learning with mathematics.           40    30    30 2.85
## 4      Mind goes blank and I am unable to think clearly when doing my math test. 60    15    25 2.50
## 10     I would like to take more math classes.      45    30    25 2.65
## 12     Math is one of my favorite subjects.         50    25    25 2.45
##
## ad
## 3  0.8127277
## 8  1.1180540
## 1  1.3533584
## 5  1.2317550
## 2  1.3484884
## 14 1.3484884
## 6  1.3337719
## 7  1.4244112
## 9  1.4104870
## 11 1.4473206
## 13 1.2258187
## 4  1.4327008
## 10 1.1821034
## 12 1.3168943
```


91/123

Likert Bar Plot



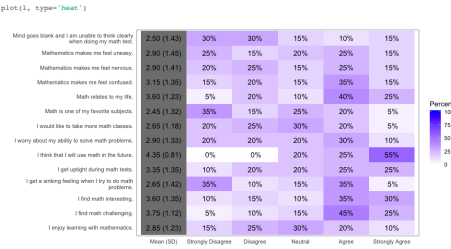
92/123

Likert Bar Plot



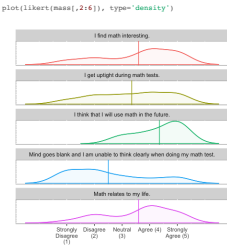
93/123

Likert Heat Plot



94/123

Likert Density Plot {flexbox vcenter}



95/123

Grouped Likert Analysis


```
lg <- Likert(masa[,2:6], grouping = masa$Gender)
summary(lg)

##      Group      Item low
## 1 Female      I find math interesting. 28.57143
## 2 Female      I get uptight during math tests. 28.57143
## 3 Female      I think that I will use math in the future. 0.00000
## 4 Female Mind goes blank and I am unable to think clearly when doing my math test. 50.00000
## 5 Female      Math relates to my life. 25.71429
## 6 Male      I find math interesting. 16.66667
## 7 Male      I get uptight during math tests. 33.33333
## 8 Male      I think that I will use math in the future. 0.00000
## 9 Male Mind goes blank and I am unable to think clearly when doing my math test. 83.33333
## 10 Male      Math relates to my life. 0.00000

##      neutral      high      mean      sd
## 1 14.28571 57.14286 3.337143 1.3924519
## 2 0.00000 71.42857 3.714286 1.3827827
## 3 28.57143 71.42857 4.285714 0.9188735
## 4 21.42857 28.57143 2.785714 1.4760288
## 5 14.28571 50.00000 3.285714 1.3259871
## 6 0.00000 83.33333 4.166667 1.1694652
## 7 66.66667 0.00000 2.500000 0.8366600
## 8 0.00000 100.00000 4.500000 0.5477224
## 9 0.00000 16.66667 1.833333 1.6894612
## 10 0.00000 100.00000 4.333333 0.5163978
```

96/123

Grouped Likert Bar Plot



97/123

Markdown

Markdown is a plain text format designed primarily to be converted to HTML documents. However, with the development of [Pandoc](#), markdown has become popular for generating lots of other document formats including PDF, Word, and presentations (this presentation is, in fact, written in markdown).

The goal of markdown is for the source file (in plain text) to be readable without converting to the desired output format.

```
# Level 1 Header
## Level 2 Header
### Level 3 Header
#### Level 4 Header
##### Level 5 Header
```

One asterisk around a word or phrase will put it in *italics*, two will make it **bold**.

Images


Links

See <https://daringfireball.net/projects/markdown/> for more on Markdown.

100/123

Pairs Plot

```
pairs(mass[,2:6], panel=panel.smooth, diag=panel=panel.hist, upper.panel=panel.cor)
```



98/123

R Markdown

R markdown extends markdown by allowing for embedding R code and output directly within your document.

```
```{r label, options}
2 x 2
```
```

R code can also be included inline by placing commands between ``` and ``` (note: not space between the opening tickmark and `x`).

A total of `` x nrow(mass)`` students responded to the survey.

A total of 20 students responded to the survey.

See <https://www.rstudio.com/wp-content/uploads/2015/02/rmarkdown-cheatsheet.pdf> for a cheat sheet on markdown and R markdown.

101/123

Document Preparation

Chunk Options

- **eval** (**TRUE** or **FALSE**) - whether the chunk is executed.
- **echo** (**TRUE** or **FALSE**) - whether the source code is copied into the final document.
- **message** (**TRUE** or **FALSE**) - whether messages are included in the final document.
- **warning** (**TRUE** or **FALSE**) - whether warning messages are included in the final document.
- **results** (**asis**, **hide**) - how the output should be included.
- **fig.width** - the width of any figures from the code chunk.
- **fig.height** - the height of any figures from the code chunk.
- **cache** - whether the R chunk should be cached and only updated if changed. Useful if the chunk takes a significant amount of time to run.

102/123


Tables

```
library(knitr)
tab <- describeby(tutoringGrade, group=tutoringtreat, mat=TRUE, show=TRUE)
tab <- tab[,c('group', 'n', 'mean', 'sd', 'median')]
names(tab) <- c('Treatment', 'n', 'mean', 'Standard Deviation', 'Median')
kable(tab, row.names=FALSE, digits=3, option="Descriptive Statistics for Grade by Treatment")
```

| Descriptive Statistics for Grade by Treatment | | | | |
|---|-----|-------|--------------------|--------|
| Treatment | n | Mean | Standard Deviation | Median |
| Control | 918 | 2.791 | 1.543 | 3 |
| Treat1 | 134 | 3.179 | 1.003 | 3 |
| Treat2 | 90 | 3.489 | 0.824 | 4 |

Math Equations

R Markdown uses Mathjax to display math equations. Mathjax is an open source JavaScript library that converts LaTeX style equations to HTML. Thanks to Daum Equation Editor available as a Google Chrome Extension or as a Mac Application, creating equations is relatively simply using a graphical user interface. Math equations are surrounded by one dollar sign for inline equations or two dollar signs for centered equations.



```
---
SS: ( * )^( \pi ) + 1 = 0 SS
---
```

$$e^x + 1 = 0$$

Affairs Data

```
library(MB)
library(party)
library(randomForest)
library(MOCH)

data(Affairs)
str(Affairs)

## 'data.frame':    601 obs. of  9 variables:
## $ affairs      : num  0 0 0 0 0 0 0 0 0 ...
## $ gender       : factor w/ 2 levels "female","male": 2 1 1 2 2 1 1 2 1 2 ...
## $ age          : num  37 27 32 57 22 22 22 57 32 22 ...
## $ yearsmarried : num  10 4 15 15 0.75 1.5 0.75 15 15 1.5 ...
## $ children     : factor w/ 2 levels "no", "yes": 1 1 2 2 1 1 1 2 2 1 ...
## $ religiousness: int   3 4 1 5 2 2 2 2 4 4 ...
## $ education    : num  18 14 12 18 17 17 12 14 16 14 ...
## $ occupation   : int   7 6 1 6 6 5 1 4 1 4 ...
## $ rating       : int   4 4 4 5 3 5 3 4 2 5 ...
```

Fair, R.C.(1978). A Theory of Extramarital Affairs. *Journal of Political Economy*, 86, 45-61.

Affairs Data (cont.)

- **affairs** - How often engaged in extramarital sexual intercourse during the past year?
- **gender** - factor indicating gender.
- **age** - numeric variable coding age in years: 17.5 = under 20, 22 = 20-24, 27 = 25-29, 32 = 30-34, 37 = 35-39, 42 = 40-44, 47 = 45-49, 52 = 50-54, 57 = 55 or over.
- **yearsmarried** - numeric variable coding number of years married: 0.125 = 3 months or less, 0.417 = 4-6 months, 0.75 = 6 months-1 year, 1.5 = 1-2 years, 4 = 3-5 years, 7 = 6-8 years, 10 = 9-11 years, 15 = 12 or more years.
- **children** - Are there children in the marriage?
- **religiousness** - numeric variable coding religiousness: 1 = anti, 2 = not at all, 3 = slightly, 4 = somewhat, 5 = very.
- **education** - numeric variable coding level of education: 9 = grade school, 12 = high school graduate, 14 = some college, 16 = college graduate, 17 = some graduate work, 18 = master's degree, 20 = Ph.D., M.D., or other advanced degree.
- **occupation** - numeric variable coding occupation according to Hollingshead classification (reverse numbering).
- **rating** - numeric variable coding self rating of marriage: 1 = very unhappy, 2 = somewhat unhappy, 3 = average, 4 = happier than average, 5 = very happy.

Advanced Statistics and Predictive Modeling

Data and Model Preparation

```
Affairs$badAffair <- Affair$affairs > 0

train <- sample(nrow(Affairs), nrow(Affairs) / 2)
affairs.train <- Affairs[train,]
affairs.valid <- Affairs[-train,]

formu <- badAffair ~ gender + age + yearsmarried + children +
  religiousness + education + occupation + rating
```

Logistic Regression

```
lr.out <- glm(formu, data=affairs.train, family=binomial())
summary(lr.out)

##
## Call:
## glm(formula = formu, family = binomial(), data = affairs.train)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -1.3925  -0.7357  -0.5349  -0.2834   2.4298
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)    2.92232    1.35559   2.158 0.028667 *
## gendermale     0.46772    0.34642   1.327 0.03921 .
## age           -0.46547    0.07727  -5.481 0.000000 ***
## yearsmarried   0.16924    0.04767   3.551 0.000384 ***
## childrenyes   -0.48637    0.41742  -1.165 0.243947
## religiousness  -0.28367    0.13195  -2.150 0.03444 *
## education     -0.03138    0.07568  -0.415 0.678352
## occupation     0.03360    0.10675   0.320 0.897153
## rating        -0.46011    0.13374  -3.440 0.000581 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
## Null deviance: 335.19 on 299 degrees of freedom
## Residual deviance: 296.05 on 291 degrees of freedom
## AIC: 314.05
##
## Number of Fisher Scoring iterations: 4

train.out <- E1tted(lr.out)
head(train.out)

##           645      1685      9012      1726       735      1573
## 0.2606440 0.2716928 0.5508259 0.1885330 0.4659040 0.1670437


valid.out <- predict(lr.out, newdata=affairs.valid, type="response")
head(valid.out)

##           5       23       44       45       50       64
## 0.1300914 0.4202230 0.2887123 0.2532270 0.7313187 0.1618934
```

109/123

Receiver Operating Characteristic (ROC) Curve

```
pred <- prediction(valid.out, affairs.valid$ndAffair)
pewf <- performance(pred, measure = "auc", x.measure = "fpr")
plot(pewf, col=rainbow(10)); abline(0, 1)
```



110/123

Confusion matrix and model accuracy

```
table(prediction=valid.out > median(valid.out), 'Actual'=affairs.valid$ndAffair) %>%
  print %>% prop.table %>% diag %>% sum %>% print

##              Actual
## Prediction FALSE TRUE
## FALSE      124    25
## TRUE       99    51
## [1] 0.5880399
```

111/123

Decision Trees

```
library(rpart)
rpart.out <- rpart(formu, data=affairs.train)
plot(rpart.out); text(rpart.out)
```



112/123

Decision Trees (cont.)

```
valid.out <- predict(rpart.out, newdata=affairs.valid)
table(prediction=valid.out > median(valid.out), 'Actual'=affairs.valid$ndAffair) %>%
  print %>% prop.table %>% diag %>% sum %>% print

##              Actual
## Prediction FALSE TRUE
## FALSE      141    28
## TRUE       84    48
## [1] 0.427907
```

113/123

Random Forests

```
form2 <- update(formu, factor(ndAffair) ~ .)
library(randomForest)
rf.out <- randomForest(form2, data=affairs.train)
valid.out.rf <- predict(rf.out, newdata=affairs.valid)
table(prediction=valid.out.rf, 'Actual'=affairs.valid$ndAffair) %>%
  print %>% prop.table %>% diag %>% sum %>% print

##              Actual
## Prediction FALSE TRUE
## FALSE      203    63
## TRUE       22    13
## [1] 0.717608
```

114/123

Shiny

[Shiny](#) is an R package that provides a framework for creating interactive, web based, applications.
[Tutorial and CheatSheet](#)

ui.R

```
shinyUI(pageWithSidebar(
  headerPanel("Gambler's Run"), # Application title
  sidebarPanel(
    helpText(paste0("Select the number of games to play (x axis), ",
      "the odds of winning, and the number of runs ",
      "(i.e. number of lines).")),
    sliderInput("games", "Number of games", min=2, max=1000, value=100),
    sliderInput("odds", "Odds of winning (1:n)", min=2, max=12, value=2),
    sliderInput("runs", "Number of runs", min=1, max=20, value=1),
    br(), actionButton("load.data", "Start Over")
  ),
  mainPanel(
    tabsetPanel(
      tabPanel("Plot", plotOutput("results")), plotOutput("plot")),
      tabPanel("Table", tableOutput("table"))
    )
  )
})
```

Interactive Documents

116/123

117/123

server.R

```
shinyServer(function(input, output) {
  data <- reactive({
    if(input$reload.data > 0 | TRUE) {
      data <- data.frame(x=integer(), y=integer(), run=integer())
      for(i in seq_len(input$games)) {
        vals <- sample(input$odds, input$games, replace=TRUE)
        vals <- ifelse(vals==1, 0, -1)
        data <- rbind(data, data.frame(x = 1:length(vals), y = cumsum(vals),
          run = rep(i, length(vals)) ))
      }
      return(data)
    }
  })
  output$table <- renderTable({ data() })
  output$plot <- renderPlot({
    mydata <- data()
    range <- c( min(abs(mydata$y)), max(abs(mydata$y)) )
    ggplot(mydata, aes(x=x, y=y, colour=run, group=run)) + geom_hline(yintercept=0, colour='blue') +
      geom_line() + ylim(range) + ylab('Cumulative Sum') + xlab('x') + theme(legend.position='none')
  }, height=400)
  output$results <- renderText({
    mydata <- data()
    total <- mean(mydata$mydata$x == max(mydata$x), 'y')
    return(paste0("Average ", ifelse(total < 0, 'losses', 'winsings'), " after ",
      input$games, " games is ", prettyNum(shot(total), digits=1)))
  })
})
```

118/123

UI Inputs

- **actionButton** - Action Button
- **checkboxGroupInput** - A group of check boxes
- **checkboxInput** - A single check box
- **dateInput** - A calendar to aid date selection
- **dateRangeInput** - A pair of calendars for selecting a date range
- **fileInput** - A file upload control wizard
- **helpText** - Help text that can be added to an input form
- **numericInput** - A field to enter numbers
- **radioButtons** - A set of radio buttons
- **selectInput** - A box with choices to select from
- **sliderInput** - A slider bar
- **submitButton** - A submit button
- **textInput** - A field to enter text

119/123

Gambler Demo

```
shiny::runGitHub('shinyApps', 'johyer', subdir='gambler')
```

120/123

Additional Resources

- [R Bloggers](#)
- [Quick-R and R in Action Book](#)
- [The Art of R Programming](#)
- Hadley Wickham's books and [websites](#): [R Packages](#), [Advanced R](#), [ggplot2](#)
- [The R Journal](#)
- [Journal of Statistical Software](#)
- RStudio [Cheat Sheets](#), [webinars](#), and [online learning](#)
- [OpenIntro](#) An open source introductory statistics textbook that make extensive use of R.
- [IS606](#) Statistics and Probability course I teach online. Most materials are available on my Github page.

Thank You!

- Jason Bryer, Ph.D. jbryer@excelsior.edu, www.bryer.org
- Kim Speerschneder kspeerschneder@excelsior.edu
- Workshop Materials: github.com/jbryer/introR

Conclusions